INDIAN DATA CLUB

CODE BASICS

databricks

14 DAYS

AI CHALLENGE

DAY 12

Topic:

MLflow Basics

Challenge:

1. Train simple regression model
2. Log parameters, metrics, model
3. View in MLflow UI
4. Compare runs

#DatabricksWithIDC

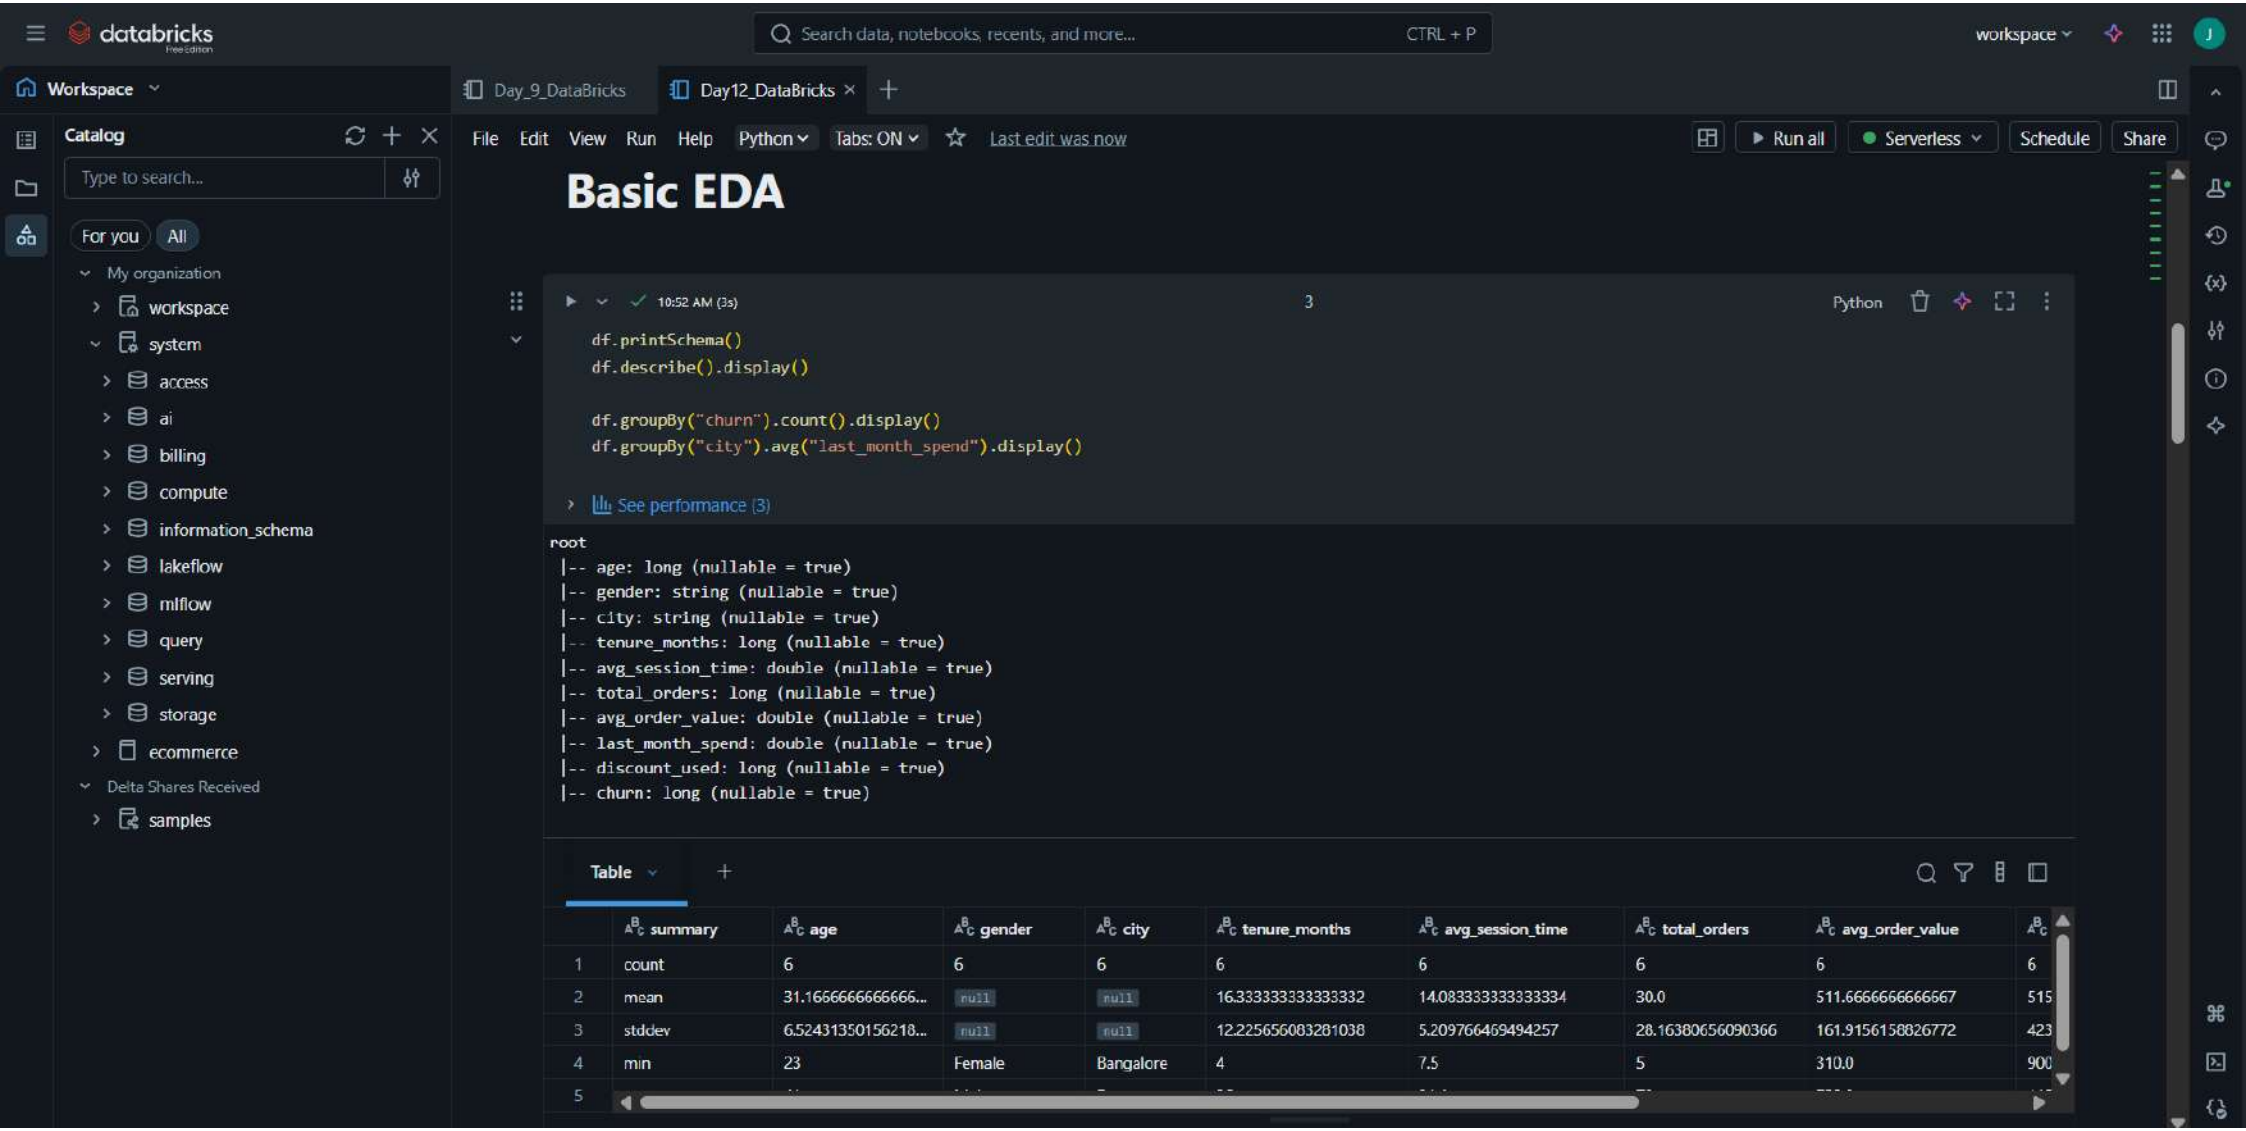
databricks
Free edition
Search data, notebooks, recents, and more...
CTRL + P
workspace
Workspace
Day_9_DataBricks
Day12_DataBricks
Catalog
Type to search...
For you
All
My organization
workspace
system
access
ai
billing
compute
information_schema
lakeflow
mlflow
query
serving
storage
ecommerce
Delta Shares Received
samples
File Edit View Run Help Python Tabs: ON Last edit was now
Run all
Serverless
Schedule
Share
Basic EDA
10:52 AM (3s)
3
Python
df.printSchema()
df.describe().display()
df.groupBy("churn").count().display()
df.groupBy("city").avg("last_month_spend").display()
See performance (3)
root
|-- age: long (nullable = true)
|-- gender: string (nullable = true)
|-- city: string (nullable = true)
|-- tenure_months: long (nullable = true)
|-- avg_session_time: double (nullable = true)
|-- total_orders: long (nullable = true)
|-- avg_order_value: double (nullable = true)
|-- last_month_spend: double (nullable = true)
|-- discount_used: long (nullable = true)
|-- churn: long (nullable = true)
Table
summary
age
gender
city
tenure_months
avg_session_time
total_orders
avg_order_value
1 count 6 6 6 6 6 6 6 6
2 mean 31.1666666666666... null null 16.333333333333332 14.083333333333334 30.0 511.6666666666667 515
3 stddev 6.52431350156218... null null 12.225656083281038 5.209766469494257 28.16380656090366 161.9156158826772 423
4 min 23 Female Bangalore 4 7.5 5 310.0 900
5

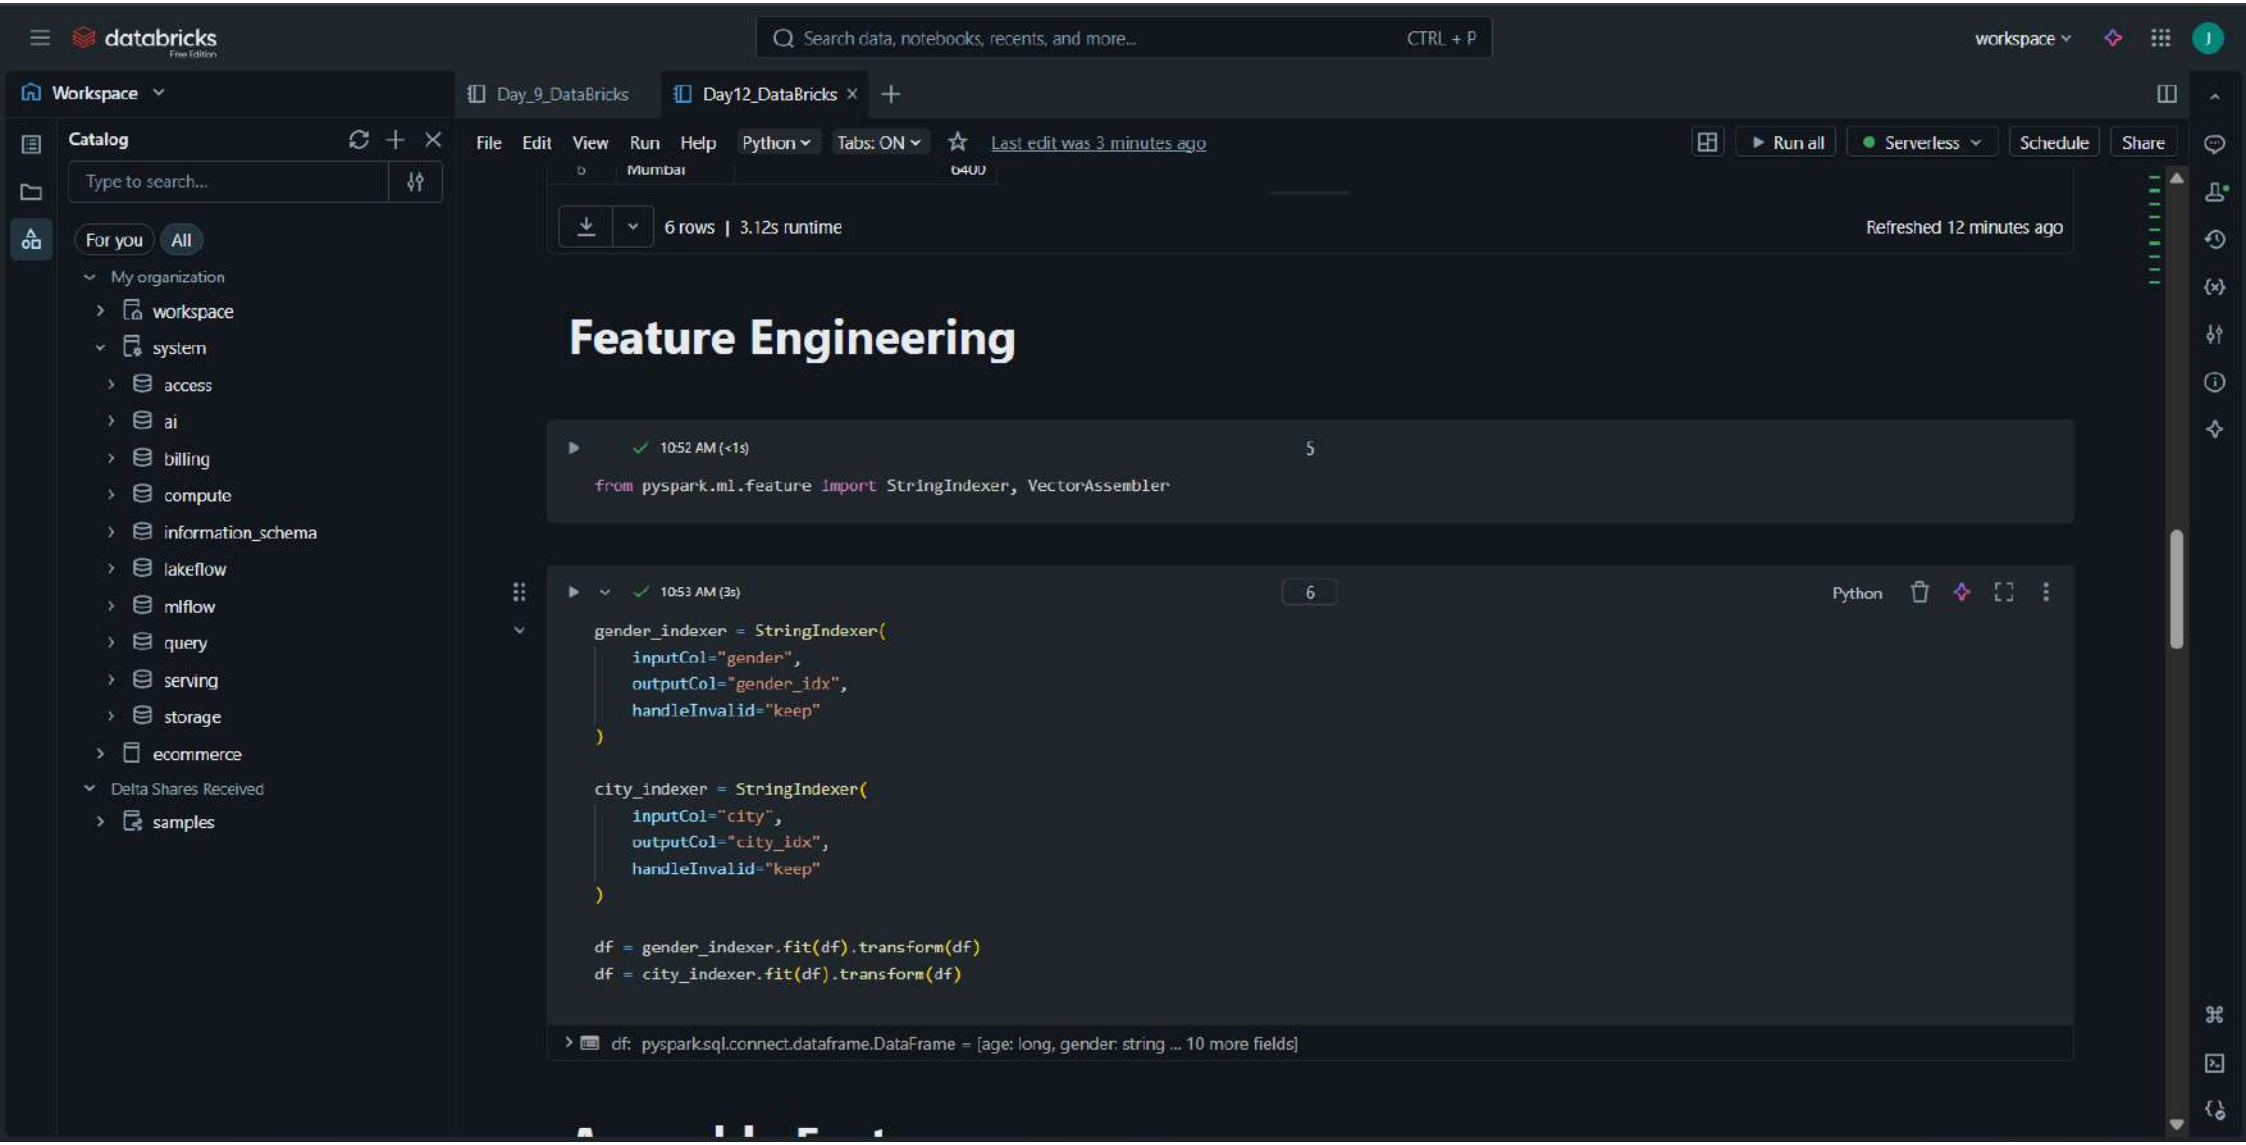
databricks
Free Edition
Search data, notebooks, recents, and more...
CTRL + P
workspace
Workspace
Day_9_DataBricks
Day12_DataBricks
Catalog
Type to search...
For you
All
My organization
workspace
system
access
ai
billing
compute
information_schema
lakeflow
mlflow
query
serving
storage
ecommerce
Delta Shares Received
samples
File Edit View Run Help
Python
Tabs: ON
Last edit was 3 minutes ago
Run all
Serverless
Schedule
Share
Mumbai 6400
6 rows | 3.12s runtime
Refreshed 12 minutes ago
Feature Engineering
10:52 AM (<1s)
5
from pyspark.ml.feature import StringIndexer, VectorAssembler
10:53 AM (3s)
6
Python
gender_indexer = StringIndexer(
inputCol="gender",
outputCol="gender_idx",
handleInvalid="keep"
)
city_indexer = StringIndexer(
inputCol="city",
outputCol="city_idx",
handleInvalid="keep"
)
df = gender_indexer.fit(df).transform(df)
df = city_indexer.fit(df).transform(df)
df: pyspark.sql.connect.dataframe.DataFrame = [age: long, gender: string ... 10 more fields]

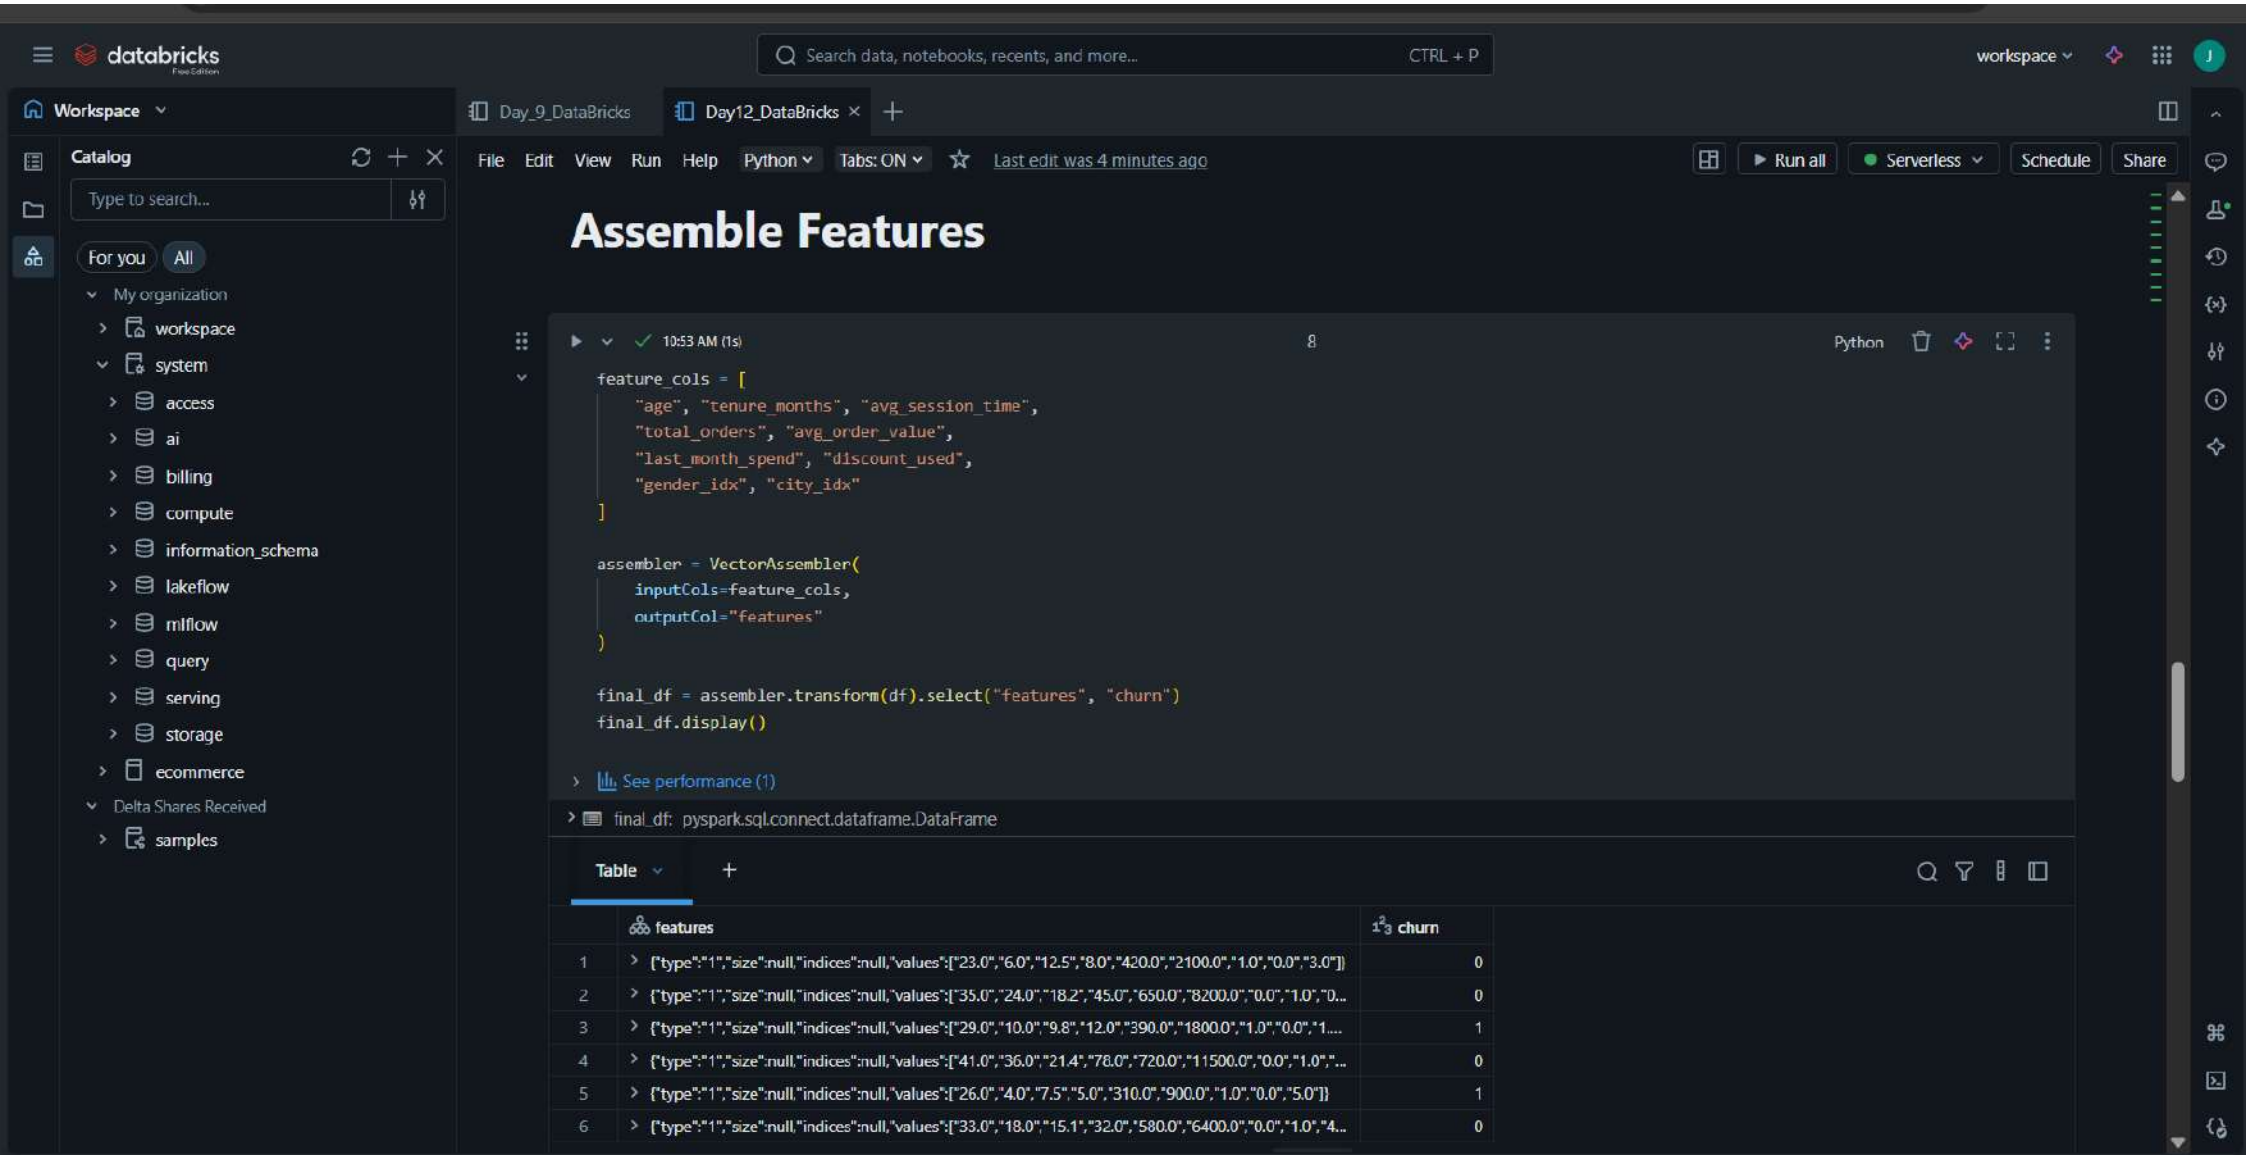
databricks
Search data, notebooks, recents, and more... CTRL + P
workspace
Workspace
Day_9_DataBricks
Day12_DataBricks
Catalog
Type to search...
For you
All
My organization
workspace
system
access
ai
billing
compute
information_schema
lakeflow
mlflow
query
serving
storage
ecommerce
Delta Shares Received
samples
File Edit View Run Help Python Tabs: ON Last edit was 4 minutes ago
Run all
Serverless
Schedule
Share
Assemble Features
10:53 AM (1s)
Python
feature_cols = [
"age", "tenure_months", "avg_session_time",
"total_orders", "avg_order_value",
"last_month_spend", "discount_used",
"gender_idx", "city_idx"
]
assembler = VectorAssembler(
inputCols=feature_cols,
outputCol="features"
)
final_df = assembler.transform(df).select("features", "churn")
final_df.display()
See performance (1)
final_df: pyspark.sql.connect.dataframe.DataFrame
Table
features
churn
1 {"type":"1","size":null,"indices":null,"values":["23.0","6.0","12.5","8.0","420.0","2100.0","1.0","0.0","3.0"]} 0
2 {"type":"1","size":null,"indices":null,"values":["35.0","24.0","18.2","45.0","650.0","8200.0","0.0","1.0","0... 0
3 {"type":"1","size":null,"indices":null,"values":["29.0","10.0","9.8","12.0","390.0","1800.0","1.0","0.0","1.... 1
4 {"type":"1","size":null,"indices":null,"values":["41.0","36.0","21.4","78.0","720.0","11500.0","0.0","1.0","... 0
5 {"type":"1","size":null,"indices":null,"values":["26.0","4.0","7.5","5.0","310.0","900.0","1.0","0.0","5.0"]} 1
6 {"type":"1","size":null,"indices":null,"values":["33.0","18.0","15.1","32.0","580.0","6400.0","0.0","1.0","4... 0

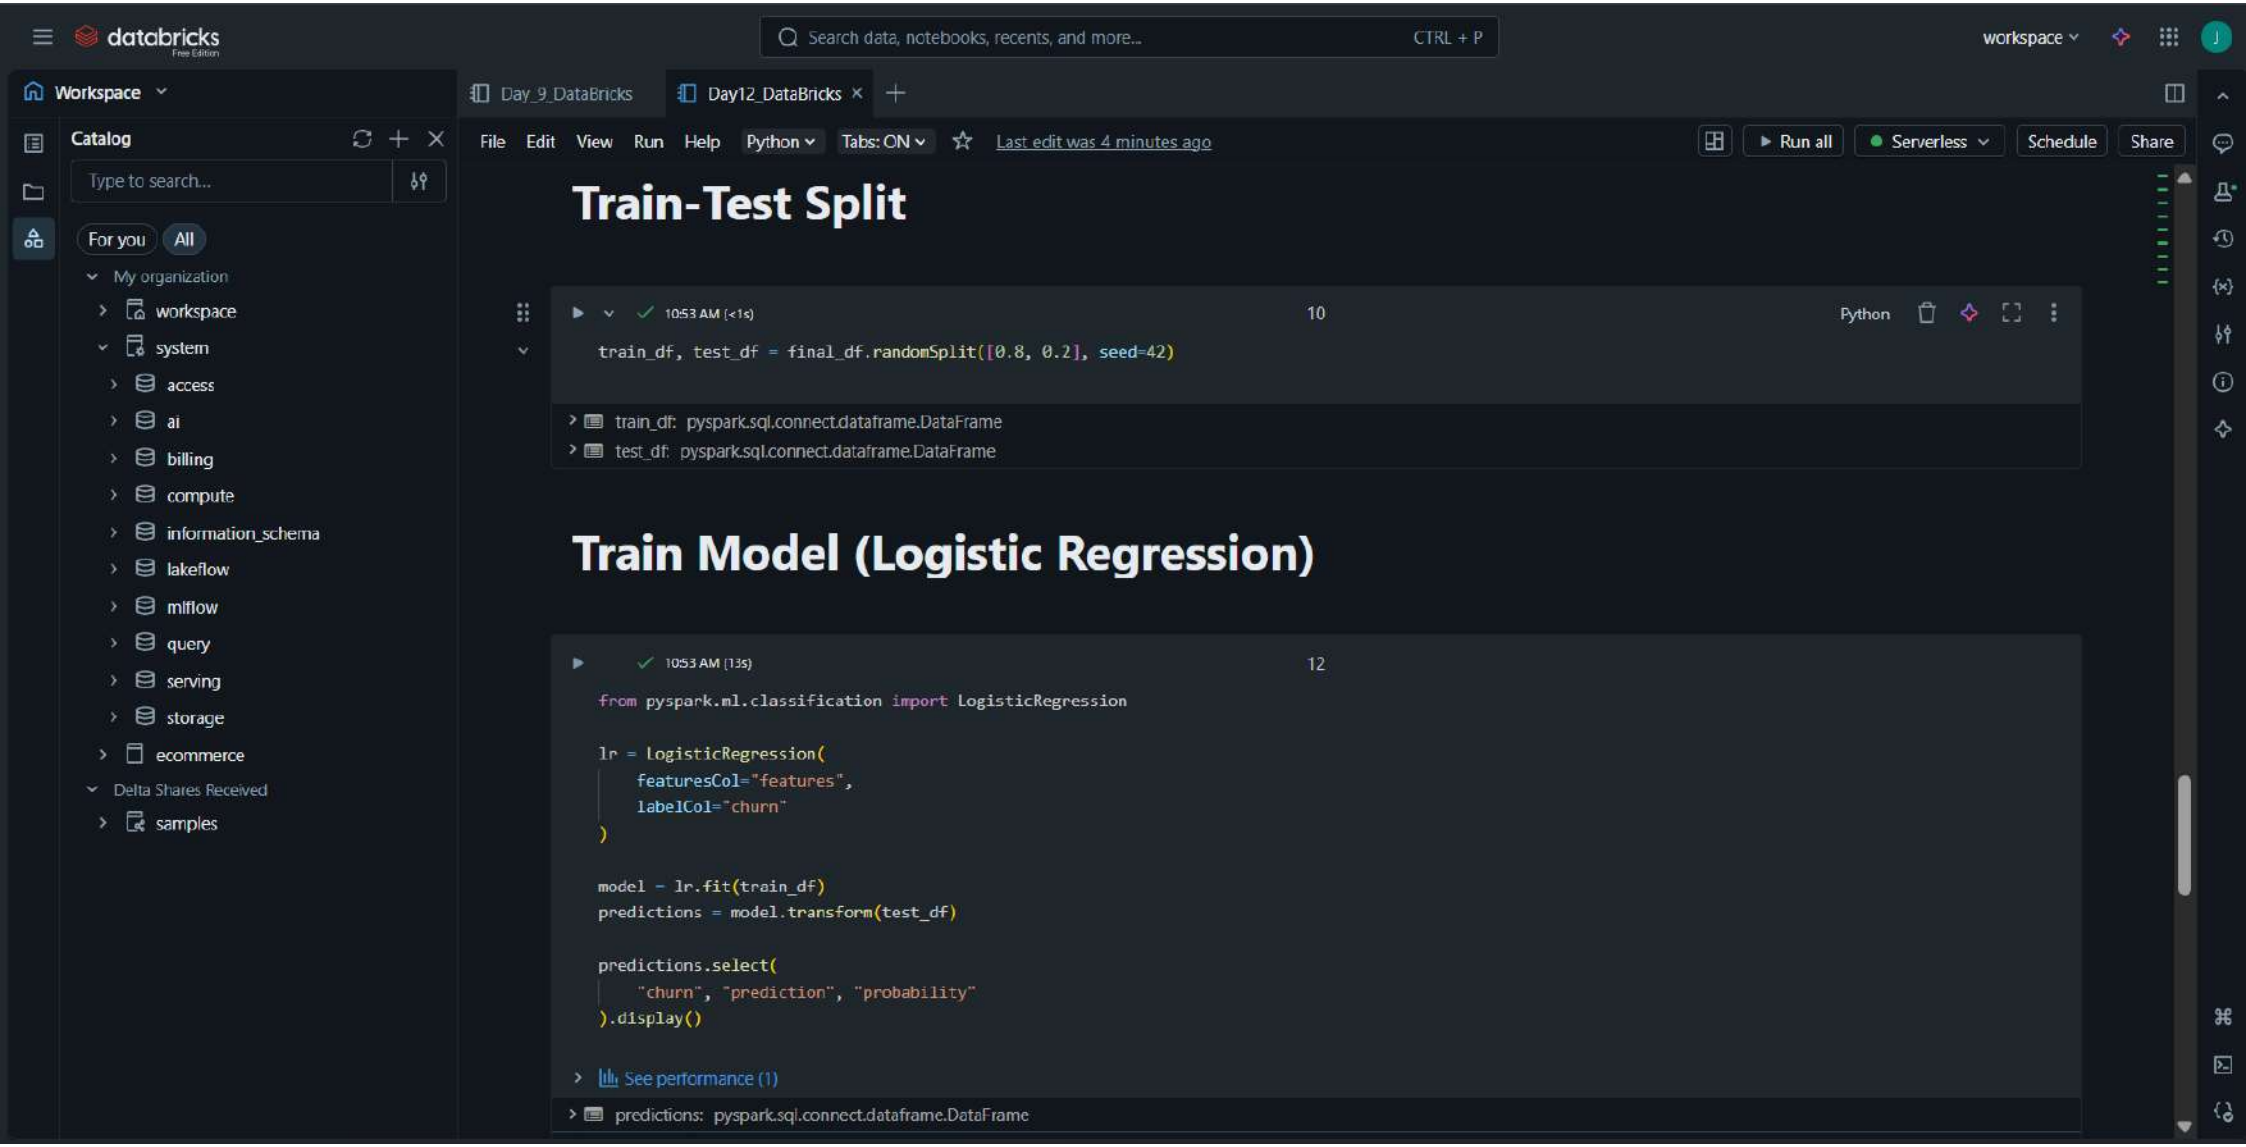
databricks
Free Edition
Search data, notebooks, recents, and more...
CTRL + P
workspace
Workspace
Day_9_DataBricks
Day12_DataBricks
Catalog
Type to search...
For you
All
My organization
workspace
system
access
ai
billing
compute
information_schema
lakeflow
mlflow
query
serving
storage
ecommerce
Delta Shares Received
samples
File
Edit
View
Run
Help
Python
Tabs: ON
Last edit was 4 minutes ago
Run all
Serverless
Schedule
Share
Train-Test Split
10:53 AM (<1s)
10
Python
train_df, test_df = final_df.randomSplit([0.8, 0.2], seed=42)
train_df: pyspark.sql.connect.dataframe.DataFrame
test_df: pyspark.sql.connect.dataframe.DataFrame
Train Model (Logistic Regression)
10:53 AM (13s)
12
from pyspark.ml.classification import LogisticRegression
lr = LogisticRegression(
featuresCol="features",
labelCol="churn"
)
model = lr.fit(train_df)
predictions = model.transform(test_df)
predictions.select(
"churn", "prediction", "probability"
).display()
See performance (1)
predictions: pyspark.sql.connect.dataframe.DataFrame

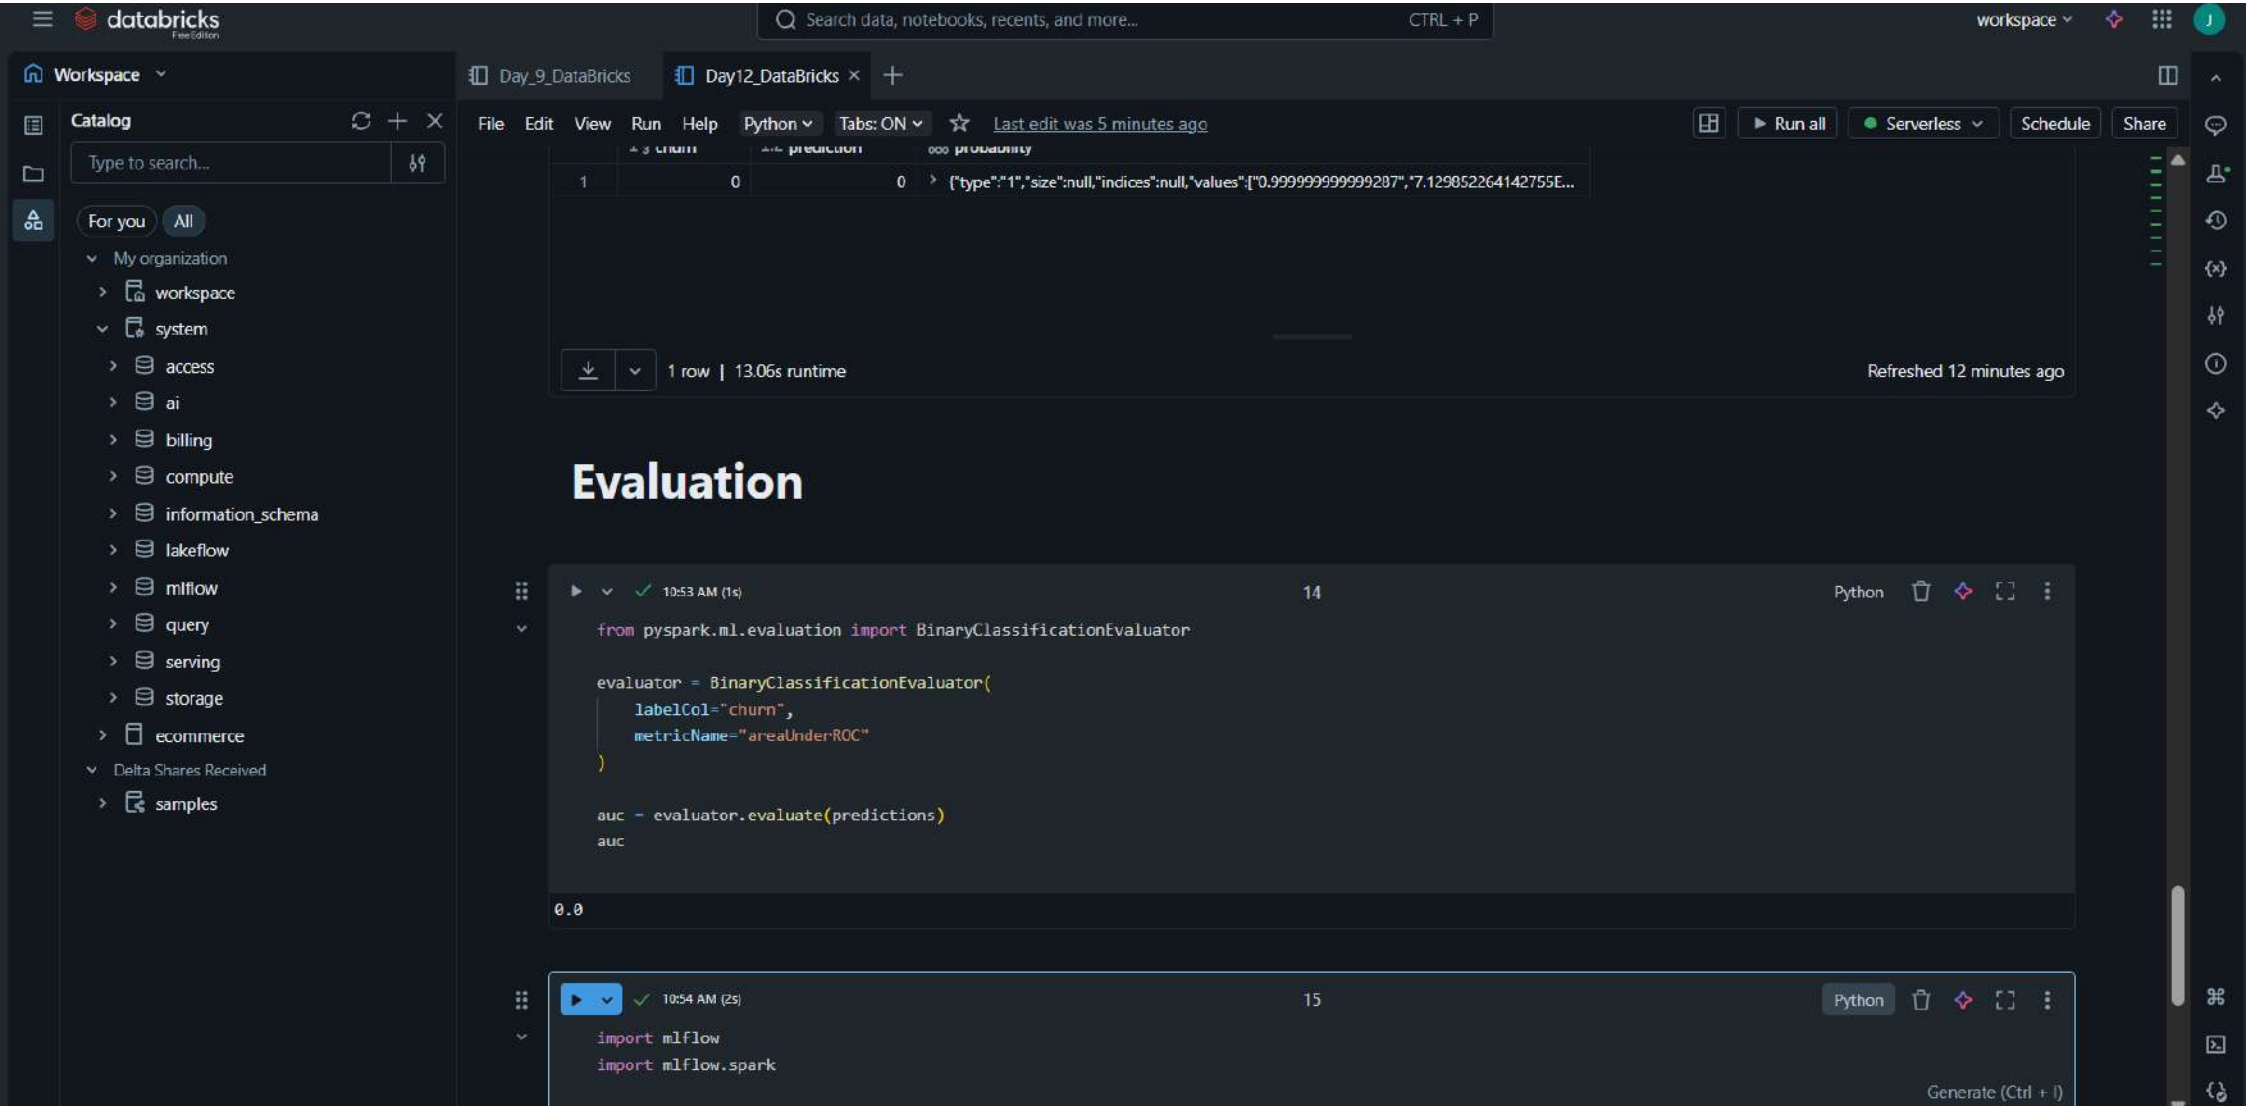
databricks
Search data, notebooks, recents, and more...
CTRL + P
workspace
Workspace
Day_9_DataBricks
Day12_DataBricks
Catalog
Type to search...
For you
All
My organization
workspace
system
access
ai
billing
compute
information_schema
lakeflow
mlflow
query
serving
storage
ecommerce
Delta Shares Received
samples
File Edit View Run Help
Python
Tabs: ON
Last edit was 5 minutes ago
Run all
Serverless
Schedule
Share
1 0 0 {"type":"1","size":null,"indices":null,"values":["0.999999999999287","7.129852264142755E...
1 row | 13.06s runtime
Refreshed 12 minutes ago
Evaluation
10:53 AM (1s)
14
Python
from pyspark.ml.evaluation import BinaryClassificationEvaluator
evaluator = BinaryClassificationEvaluator(
labelCol="churn",
metricName="areaUnderROC"
)
auc = evaluator.evaluate(predictions)
auc
0.0
10:54 AM (2s)
15
Python
import mlflow
import mlflow.spark
Generate (Ctrl + I)