



14 DAYS

AI CHALLENGE

DAY 14

Topic:

AI-Powered Analytics: Genie & Mosaic AI

Challenge:

1. Use Genie to query data with natural language
2. Explore Mosaic AI features
3. Build simple NLP task
4. Create AI-powered insights

Databricks Free Edition

Search data, notebooks, recents, and more... CTRL + P

workspace ▾

Workspace Catalog

Type to search...

For you All

- My organization
- workspace
- system
- ecommerce
- Delta Shares Received
- samples

Day13_DatabRicks Day14_DataBrics +

File Edit View Run Help Python Tabs: ON Last edit was 6 hours ago

Run all Serverless Schedule Share

01:28 PM (37s) 1

```
%pip install torch transformers
```

Collecting torch
 Downloading torch-2.10.0-cp312-cp312-manylinux_2_28_aarch64.whl.metadata (31 kB)
Collecting transformers
 Downloading transformers-4.57.6-py3-none-any.whl.metadata (43 kB)
Requirement already satisfied: filelock in /usr/local/lib/python3.12/dist-packages (from torch) (3.18.0)
Requirement already satisfied: typing-extensions>=4.10.0 in /databricks/python3/lib/python3.12/site-packages (from torch) (4.12.2)
Requirement already satisfied: setuptools in /usr/local/lib/python3.12/dist-packages (from torch) (74.0.0)
Collecting sympy>=1.13.3 (from torch)
 Downloading sympy-1.14.0-py3-none-any.whl.metadata (12 kB)
Collecting networkx>=2.5.1 (from torch)
 Downloading networkx-3.6.1-py3-none-any.whl.metadata (6.8 kB)
Requirement already satisfied: jinja2 in /databricks/python3/lib/python3.12/site-packages (from torch) (3.1.5)
Requirement already satisfied: fsspec>=0.8.5 in /databricks/python3/lib/python3.12/site-packages (from torch) (2023.5.0)
Collecting huggingface-hub<1.0,>=0.34.0 (from transformers)
 Downloading huggingface_hub-0.36.0-py3-none-any.whl.metadata (14 kB)
Requirement already satisfied: numpy>=1.17 in /databricks/python3/lib/python3.12/site-packages (from transformers) (2.1.3)
Requirement already satisfied: packaging>=20.0 in /databricks/python3/lib/python3.12/site-packages (from transformers) (24.1)
Requirement already satisfied: pyyaml>=5.1 in /databricks/python3/lib/python3.12/site-packages (from transformers) (6.0.2)
Collecting regex!=2019.12.17 (from transformers)
 Downloading regex-2026.1.15-cp312-cp312-manylinux2014_aarch64_manylinux_2_17_aarch64_manylinux_2_28_aarch64.whl.metadata (40 kB)
Requirement already satisfied: requests in /databricks/python3/lib/python3.12/site-packages (from transformers) (2.32.3)

💡 1

01:33 PM (7s) 2

```
from transformers import pipeline  
import mlflow
```

/local_disk0/_ephemeral_nfs/envs/pythonEnv-e1fc4b6-a796-4390-9c98-fbac2194d501/lib/python3.12/site-packages/torch/_vman_internals.py:9: FutureWarning: `isinstance`

databricks
Free Edition

Search data, notebooks, recents, and more... CTRL + P

workspace ▾ J

Workspace Catalog +

File Edit View Run Help Python Tabs: ON Last edit was 6 hours ago

For you All

- My organization
 - workspace
 - system
 - ecommerce
- Delta Shares Received
- samples

Day14.Databricks +

3 minutes ago (12s) 2

```
from transformers import pipeline
import mlflow
```

/local_disk0/.ephemeral_nfs/envs/pythonEnv-75ca6bb9-5ef6-4a05-aaed-ad536e325529/lib/python3.12/site-packages/torch/_vmap_internals.py:9: FutureWarning: `isinstance(treespec, LeafSpec)` is deprecated, use `isinstance(treespec, TreeSpec)` and `treespec.is_leaf()`` instead.
from torch.utils._pytree import _broadcast_to_and_flatten, tree_flatten, tree_unflatten

💡 1

Build a simple NLP task

2 minutes ago (2s) 4

```
# Load sentiment analysis pipeline (Mosaic AI / HuggingFace compatible)
classifier = pipeline("sentiment-analysis")
```

No model was supplied, defaulted to distilbert/distilbert-base-uncased-finetuned-sst-2-english and revision 714eb0f (<https://huggingface.co/distilbert/distilbert-base-uncased-finetuned-sst-2-english>).
Using a pipeline without specifying a model name and revision in production is not recommended.

config.json: 100% [██████████] 629/629 [00:00<00:00, 53.5kB/s]

model.safetensors: 100% [██████████] 268M/268M [00:00<00:00, 297MB/s]

tokenizer_config.json: 100% [██████████] 48.0/48.0 [00:00<00:00, 9.02kB/s]

vocab.txt: [██████████] 232k? [00:00<00:00, 26.5MB/s]

Device set to use cpu

databricks
Feed Edition

Search data, notebooks, recents, and more... CTRL + P

workspace ▾

Workspace Catalog +

File Edit View Run Help Python Tabs: ON Last edit was 6 hours ago

For you All

- My organization
 - workspace
 - system
 - ecommerce
- Delta Shares Received
- samples

Sample product reviews

```
reviews = [
    "This product is amazing and works perfectly!",
    "Terrible quality, totally disappointed",
    "Good value for money",
    "Worst purchase I have ever made"
]

results = classifier(reviews)
results
```

[{'label': 'POSITIVE', 'score': 0.9998794794082642}, {'label': 'NEGATIVE', 'score': 0.9997342228889465}, {'label': 'POSITIVE', 'score': 0.9998501539230347}, {'label': 'NEGATIVE', 'score': 0.9997796416282654}]

Create AI-powered insights

4 minutes ago (10s)

```
sentiment_df = spark.createDataFrame(
    [(reviews[i], results[i]['label'], results[i]['score']) for i in range(len(reviews))],
    ["review_text", "sentiment", "confidence"]
)

sentiment_df.display()
```

See performance ()

Screenshot of the Databricks workspace interface. The left sidebar shows the Catalog and workspace navigation. The main area displays a DataFrame named `sentiment_df` with the following data:

	review_text	sentiment	confidence
1	This product is amazing and works perfect...	POSITIVE	0.9998794794082642
2	Terrible quality, totally disappointed	NEGATIVE	0.9997342228889455
3	Good value for money	POSITIVE	0.9998501539230347
4	Worst purchase I have ever made	NEGATIVE	0.9997796416282654

The DataFrame was last edited 6 hours ago and was refreshed 3 minutes ago.

Mosaic AI Exploration (MLflow logging)

Screenshot of a Databricks notebook cell containing Python code for MLflow logging:

```
with mlflow.start_run(run_name="sentiment_analysis_model"):
    mlflow.log_param("model", "distilbert-base-uncased-finetuned-sst-2")
    mlflow.log_metric("avg_confidence",
                      sum([r["score"] for r in results]) / len(results))
```

The cell output shows:

- (1) MLflow run
- Logged 1 run to an experiment in MLflow. Learn more

At the bottom of the cell, keyboard shortcuts are listed:

- [Shift+Enter] to run and move to next cell
- [Ctrl+Shift+P] to open the command palette
- [Esc H] to see all keyboard shortcuts